

BA2: Digital Korea

Week 9: Sentiment Analysis

Steven Denney

Korean Studies
Leiden University

April 13, 2026

Today's Agenda

1. Week 8 review & check-in
2. What is sentiment analysis?
3. Dictionary-based approaches
4. Sentiment arithmetic: word + word + word = ?
5. Korean sentiment dictionaries
6. Application: Moon Jae-in's tweets
7. Demo: Sentiment analysis in Orange Data Mining
8. Looking ahead

Review & Check-In

Week 8 Review

Last week we covered:

- From sparse vectors (BoW) to dense vectors (embeddings)
- Word embedding concepts and nearest neighbors
- Semantic search and word similarity
- Document embeddings: representing entire documents as vectors

Key takeaway from Week 8

Embeddings capture **meaning**: words that appear in similar contexts end up near each other in vector space.

경제 - 성장 + 교육 $\approx$ 발전 (vector arithmetic)

Week 8 Review: Word Embeddings

The big idea: words as vectors capture semantic relationships.

$$\boxed{\text{한국}} - \boxed{\text{서울}} + \boxed{\text{도쿄}} \approx \boxed{\text{일본}}$$

Today we use a similar intuition, but instead of **meaning**, we measure **feeling**.

From meaning to sentiment

Word embeddings ask: *what does this word mean?*

Sentiment analysis asks: *how does this word feel?*

What Is Sentiment Analysis?

What Is Sentiment Analysis?

Sentiment analysis = determining whether a text expresses a positive, negative, or neutral attitude.

Examples in everyday life

- Product reviews (5 stars?)
- Social media posts (happy? angry?)
- News coverage (favorable? critical?)
- Political speeches (hopeful? alarming?)

Why it matters for research

- Track public opinion over time
- Compare tone across leaders
- Detect emotional framing
- Measure media bias

In computational text analysis

Sentiment analysis is a form of **measurement**: we assign a numeric value (valence) to text that captures its emotional direction.

Two Main Approaches

Dictionary-based (this week)

- Humans create word lists
- Each word gets a sentiment score
- Computer looks up and sums scores
- Transparent and interpretable
- No training data needed

Also called: lexicon-based, rule-based

Machine learning-based


- Train a model on labeled examples
- Model learns patterns from data
- Can capture context and nuance
- Requires training data
- Less transparent (“black box”)

Also called: supervised classification

This week's focus

Dictionary-based methods — where the researcher defines the rules and the computer applies them at scale.

How Dictionary-Based Sentiment Works



The logic: look up each word's score (-2 to +2), then sum them.

A Simple Example

Input: 경제가 성장하고 국민이 행복합니다

(The economy is growing and the people are happy)

Kiwi stems: 경제 성장 국민 행복

Stem	KNU score	Meaning
경제 (economy)	0	not in dictionary
성장 (growth)	+1	positive
국민 (people)	0	not in dictionary
행복 (happiness)	+2	very positive
Total	+3	

Result

행복 scores +2 (strong) while 성장 scores +1 (mild) — the dictionary captures **intensity**, not just polarity.

Sentiment Arithmetic

Sentiment Arithmetic: Word + Word + Word = ?

Last week: **word** + **word** = **meaning** (embedding analogies)

This week: **word** + **word** + **word** = **feeling** (sentiment scoring)



Each word contributes a score. The sum determines the overall sentiment.

Your Turn: Guess the Sentiment

What is the sentiment of each combination?

1. 평화 (peace) + 협력 (cooperation) + 대화 (dialogue) = ?
2. 전쟁 (war) + 위협 (threat) + 공포 (fear) = ?
3. 성장 (growth) + 위기 (crisis) + 극복 (overcome) = ?
4. 개혁 (reform) + 실패 (failure) + 희망 (hope) = ?

Think about it

Cases 3 and 4 are **mixed**. This is where dictionary methods get interesting — and where their limitations start to show.

Sentiment Arithmetic: Answers

#	Words	Scores	Result
1	평화 + 협력 + 대화	$(+2) + (+1) + (+1)$	Positive (+4)
2	전쟁 + 위협 + 공포	$(-2) + (-2) + (-2)$	Negative (-6)
3	성장 + 위기 + 극복	$(+1) + (-1) + (+2)$	Positive (+2)
4	개혁 + 실패 + 희망	$(+1) + (-2) + (+2)$	Positive (+1)

Key insight

Dictionary methods treat every word independently. They don't understand *context*: “overcome a crisis” and “cause a crisis” would score the same for 위기.

When Arithmetic Breaks Down

Dictionary-based sentiment has known limitations

1. **Negation:** 행복하지 않다 (not happy) — “happy” is positive, but the sentence is negative
2. **Sarcasm:** 정말 잘하셨네요 (you did really well) — could be sincere or sarcastic
3. **Domain specificity:** 급등 (surge) — positive for stocks, negative for infections
4. **Mixed sentiment:** 어렵지만 희망적이다 (difficult but hopeful) — both negative and positive

So why use dictionaries?

Despite these limits, dictionary methods are **transparent**, **reproducible**, require **no training data**, and work well at scale when you understand their assumptions.

Korean Sentiment Dictionaries

Korean Sentiment Dictionaries: The Landscape

Several research dictionaries are publicly available for Korean

Dictionary	Size	Scale	Notes
Chen & Skiena (2014)	44 langs	binary	Orange's built-in (what we use)
KNU (Park et al. 2018)	14,854	-2 to +2	Intensity scoring; research standard
KOSAC (Jang et al. 2013)	~11,275	5-class + intensity	SNU manual annotation on news
NRC Emotion Lexicon KSenticNet	~14,000 5,465	8 emotions + $\pm$ continuous	Cross-lingual, MT Korean Neural / concept-level

All legitimate, each with trade-offs

We use **Chen & Skiena** because Orange ships it. KNU is in the repo for reference if you continue with Korean sentiment.

One Method, Two Places to See It

Same dictionary, same formula — in Orange and in the interactive

In Orange

- File → Corpus
- Sentiment Analysis (Multilingual, Korean)
- Box Plot

Four widgets. No Python. Orange handles the tokenization and lookup.

In the interactive

- Pre-computed using the same Chen & Skiena lexicon
- Same scoring formula as Orange
- Step through 6 panels to see the method at work

Your Orange Box Plot matches the interactive's distribution.

Why both?

Orange gives you the tool. The interactive opens it up so you can see what the tool is doing to your tweets.

The KNU Korean Sentiment Lexicon

An alternative dictionary with intensity scoring — available in the course repo

Score	Meaning	Example	Translation
+2	Very positive	감사, 행복, 희망	gratitude, happiness, hope
+1	Positive	평화, 발전, 극복	peace, development, overcome
0	Neutral	—	—
-1	Negative	걱정, 위기, 힘들	worry, crisis, hard
-2	Very negative	가난, 파괴, 화재	poverty, destruction, fire

Why KNU might be worth using

The intensity scale (−2 to +2) captures *strength* of sentiment, not just direction. 감사 (+2) is stronger than 평화 (+1). 파괴 (−2) is stronger than 걱정 (−1). For this course we use Chen & Skiena (Orange's built-in) for simplicity — KNU is in the repo as a reference if you continue working with Korean sentiment.

What's in the KNU Dictionary?

Content-word stems — same form Kiwi produces from tweets

Positive stems (samples)

- 감사 (gratitude) → +2
- 행복 (happiness) → +2
- 희망 (hope) → +2
- 축하 (celebration) → +2
- 발전 (development) → +1
- 평화 (peace) → +1

Negative stems (samples)

- 가난 (poverty) → -2
- 파괴 (destruction) → -2
- 전쟁 (war) → -2
- 화재 (fire disaster) → -2
- 걱정 (worry) → -1
- 위기 (crisis) → -1

Available in the course repo

KNU ships in `data/sentiment_dic/` — `SentiWord_Dict.txt` (full, -2 to +2) and stem-indexed variants. Load into Orange via the Custom Dictionary option, or into R with `read_tsv()`. For Week 9 we use Orange's built-in Multilingual method instead.

Building Custom Dictionaries

Sometimes general dictionaries miss domain-specific sentiment

1. **Start with a general dictionary** (our positive/negative lists)
2. **Read a sample of your texts** — note domain-specific words
3. **Add domain terms:**
 - Political: 민주화 (democratization) → positive? 독재 (dictatorship) → negative?
 - Economic: 호황 (boom) → positive. 불황 (recession) → negative
4. **Validate:** check a sample of scored documents — do the scores make sense?

Remember GRS Chapter 2

Dictionary construction requires **domain knowledge**. You are making analytical choices about what counts as positive or negative in your research context.

Application: Moon Jae-in's Tweets

Case Study: Moon Jae-in on Twitter

3,148 tweets from @moonriver365 (2012–2020), scored with Orange's Multilingual method

The corpus: @moonriver365

- 3,148 tweets (2012–2020), 9 columns
- Metadata: date, favorites, retweets
- Three periods: Pre-presidency (1,973), Transition (393), Presidency (782)

Moon Jae-in (문재인) — 19th President (2017–2022)

- Key topics: inter-Korean relations, COVID-19, Japan trade dispute, democracy
- Active Twitter presence, especially 2012 campaign (1,427 tweets)
- Tweeting declined sharply after taking office

Research questions:

1. Is the overall tone positive or negative?
2. Does sentiment change across the three periods?
3. Do different topics carry different sentiment?
4. How does sentiment relate to real-world events?

Sentiment Arithmetic with Real Tweets

Illustrated with KNU scores — Orange uses a different dictionary, same method

Tweet (2018-10-17, presidency):

평화를 염원하는 우리 국민에게 큰 힘이 될 것입니다. 우리는 기필코 분단을 극복해낼 것입니다.

(It will be a great strength for our people who wish for peace. We will surely overcome division.)

Korean word	KNU score	
평화 (peace)	+1	positive
극복 (overcome)	+1	positive
Total	+2	

Not every meaningful word appears in a sentiment dictionary. 분단 (division) is here in the tweet but not in KNU — a known limit. Orange's Multilingual method uses a different word list so your numbers will differ, but the logic is identical.

Sentiment Arithmetic with Real Tweets

Mixed-sentiment tweets are where dictionary methods get interesting

Tweet (2020-04-29, presidency):

‘코로나19’를 극복하고, 이천 화재의 슬픔을 이겨내며, 반드시 우리의 새로운 일상을 만들어내겠습니다.


(We will overcome COVID-19, endure the grief of the Icheon fire, and build our new normal.)

Korean word	KNU score	
극복 (overcome)	+1	positive
화재 (fire disaster)	-2	very negative
슬픔 (grief)	-2	very negative
이기 (win, endure)	+1	positive
Total	-2	

Mildly negative (-2) with KNU — very-negative words outweigh positives. A human reader sees *resolve* to overcome. Orange’s Multilingual is binary, so it scores this differently. The dictionary choice matters.

What Can Sentiment Over Time Reveal?

Our corpus spans three distinct political periods



- Does sentiment differ across the three periods? (Your assignment!)
- Do real-world events (summits, trade dispute, COVID) produce visible shifts?
- Political communication tends toward a positive baseline — does it hold?

Moment in Time: The Impeachment Vote

Tweet (2016-12-09, transition, 4,445 favorites):

국민이 이겼습니다. 대통령 탄핵은 끝이 아니라 시작입니다.

(The people have won. The presidential impeachment is not the end but the beginning.)

Kiwi + KNU scoring:

- 이기 (win) → +1
- 탄핵 (impeachment) not in KNU
- Score: +1

Context:

- Posted day after the impeachment vote
- Tone is triumphant
- Even +1 understates it — the dictionary can't measure political significance

Moment in Time: COVID Zero

Tweet (2020-04-30, presidency, 25,952 favorites — most-liked COVID tweet):

72일 만의 코로나19 국내 확진자 0명. 총선 이후 14일간 선거로 인한 감염 0명. 대한민국의 힘, 국민의 힘입니다.

(For the first time in 72 days, zero domestic COVID-19 cases. Zero election-related infections in 14 days since the general election. This is the strength of Korea, the strength of the people.)

The dictionary problem

확진자 (confirmed cases) and 감염 (infection) are negative words.

But both appear next to **0명** (zero people) — the context is celebration, not alarm.

A dictionary cannot read “zero infections” as good news. A human can.

Critical Reading: What the Dictionary Misses

Tweet (2019-08-02, presidency, 19,110 favorites):

우리는 다시는 일본에게 지지 않을 것입니다. 적지 않은 어려움이 예상되지만, 우리 기업들과 국민들에겐 그 어려움을 극복할 역량이 있습니다.

(We will never again lose to Japan. Considerable difficulties are expected, but our companies and people have the capability to overcome them.)

Kiwi + KNU scoring:

- 어려움 (difficulty) → -2
- 예상 (expect) → -1
- 극복 (overcome) → +1
- Score: -2 (negative)

Human reading:

- Tone is defiant and assertive
- Negative *toward Japan* specifically
- Positive *self-framing* (national strength)
- “-2” misses the defiance

Lesson

Dictionary-based sentiment measures **word-level valence**. It doesn't capture the **target** (negative toward whom?) or the **rhetorical purpose** (rallying, not despairing).

Orange Data Mining Demo

Preprocessing for Sentiment

Orange's Multilingual method handles this for you — no code required

The widget does three things:

1. Tokenize each tweet on whitespace and punctuation
2. Look up each unique token in the positive and negative word lists
3. Score: $100 \times (\text{pos hits} - \text{neg hits}) / \text{total tokens}$

No stopwords, no stemming, no POS filter

Dictionary-based sentiment works on surface forms. Chen & Skiena's Korean list contains words like 좋은, 감사합니다, 같은 in the forms they appear in text — Orange handles the matching for you.
(Advanced path: Kiwi morphological tokenization + KNU intensity dictionary, see `data/scripts/.`)

Demo: Sentiment Analysis in Orange

Four widgets, no Python, built-in Korean support

1. Load the data

- **File** → select `moon_twitter.csv`
- Pass to **Corpus** → set text feature to text

2. Score sentiment


- **Sentiment Analysis** widget
- Method: **Multilingual Sentiment**
- Language: **Korean**
- Adds a sentiment column to the corpus

3. Explore

- **Box Plot** → subgroup by period3
- **Corpus Viewer** (optional) → sort by sentiment

Orange Workflow Overview

Four widgets in a line — no Python, no custom dictionary



What each widget does

- **File** — reads the tweet CSV into Orange
- **Corpus** — marks text as the document body
- **Sentiment Analysis** (Multilingual, Korean) — tokenizes, looks up each token in Chen & Skiena's Korean lexicon, outputs a sentiment column
- **Box Plot** — compare sentiment distributions across period3

Using the Sentiment Analysis Widget

The Multilingual method supports Korean out of the box

Settings to choose:

1. Drop a **Sentiment Analysis** widget onto the canvas
2. **Method** → “Multilingual Sentiment”
3. **Language** → “Korean”
4. Connect the output to a **Box Plot** or **Data Table**

What's happening inside

Chen & Skiena's (2014) Korean positive and negative word lists. For each tweet: count unique matches, output $100 \times (\text{pos} - \text{neg}) / \text{num tokens}$. The interactive exercise shows this in action. **Custom Dictionary** (advanced) lets you swap in your own word lists like KNU.

Interpreting Your Results

Once you have sentiment scores, ask:

1. **Distribution:** What is the overall shape? Mostly positive? Skewed?
2. **Outliers:** Which documents scored extremely high or low? Do they make sense when you read them?
3. **Patterns:** Does sentiment correlate with time, topic, or another variable?
4. **Validation:** Read 10–20 documents across the score range. Does the algorithm's judgment match yours?

Always validate!

Grimmer, Roberts & Stewart: “All quantitative models of language are wrong — but some are useful” (Ch. 2). Validation is not optional.

Looking Ahead

Week 10: Topic Modeling (LDA)

- Latent Dirichlet Allocation (LDA)
- Discovering hidden thematic structure in a corpus
- Choosing the number of topics
- Interpreting and labeling topics

Recommended reading:

- Grimmer, Roberts & Stewart — Chapter 13: Topic Models

For Next Week

Interactive exercise, assignment, and preparation

Interactive exercise:

- Explore the **Sentiment Analysis: Moon Jae-in's Tweets** interactive on the course website
- Step through all 6 panels to see exactly what Orange is doing with your tweets

Assignment:

- Load `moon_twitter.csv` into Orange
- Connect **Sentiment Analysis** (Multilingual, Korean) → **Box Plot** (subgroup: `period3`)
- Write 2–3 sentences: which period is most positive, and which word(s) in the interactive help you understand why?

Bonus challenge:

- Load the KNU dictionary from `data/sentiment_dic/` via Custom Dictionary. How do the results compare?